

In [66]: `pip install --upgrade pip`

Requirement already satisfied: pip in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (24.3.1)
Note: you may need to restart the kernel to use updated packages.

In [67]: `pip list | grep -E 'tensorflow|keras'`

```
keras                3.6.0
keras-core           0.1.7
keras-hub            0.18.0
keras-nlp            0.18.0
tensorflow           2.18.0
tensorflow-io-gcs-filesystem 0.37.1
tensorflow-text      2.18.0
```

Note: you may need to restart the kernel to use updated packages.

In [68]: `!pip install keras-core --upgrade`
`!pip install -q keras-nlp --upgrade`

Requirement already satisfied: keras-core in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (0.1.7)
Requirement already satisfied: absl-py in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from keras-core) (2.1.0)
Requirement already satisfied: numpy in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from keras-core) (2.0.2)
Requirement already satisfied: rich in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from keras-core) (13.9.4)
Requirement already satisfied: namex in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from keras-core) (0.0.8)
Requirement already satisfied: h5py in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from keras-core) (3.12.1)
Requirement already satisfied: dm-tree in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from keras-core) (0.1.8)
Requirement already satisfied: markdown-it-py>=2.2.0 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from rich->keras-core) (3.0.0)
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from rich->keras-core) (2.18.0)
Requirement already satisfied: typing-extensions<5.0,>=4.0.0 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from rich->keras-core) (4.12.2)
Requirement already satisfied: mdurl~=0.1 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from markdown-it-py>=2.2.0->rich->keras-core) (0.1.2)

In [69]: `!pip install tensorflow tensorflow-text`
`!pip install wordcloud`

Requirement already satisfied: tensorflow in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (2.18.0)

Requirement already satisfied: tensorflow-text in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (2.18.0)

Requirement already satisfied: absl-py>=1.0.0 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorflow) (2.1.0)

Requirement already satisfied: astunparse>=1.6.0 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorflow) (1.6.3)

Requirement already satisfied: flatbuffers>=24.3.25 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorflow) (24.3.25)

Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>=0.2.1 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorflow) (0.6.0)

Requirement already satisfied: google-pasta>=0.1.1 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorflow) (0.2.0)

Requirement already satisfied: libclang>=13.0.0 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorflow) (18.1.1)

Requirement already satisfied: opt-einsum>=2.3.2 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorflow) (3.4.0)

Requirement already satisfied: packaging in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorflow) (24.2)

Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!4.21.3,!4.21.4,!4.21.5,<6.0.0dev,>=3.20.3 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorflow) (5.28.3)

Requirement already satisfied: requests<3,>=2.21.0 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorflow) (2.32.3)

Requirement already satisfied: setuptools in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorflow) (65.5.0)

Requirement already satisfied: six>=1.12.0 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorflow) (1.16.0)

Requirement already satisfied: termcolor>=1.1.0 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorflow) (2.5.0)

Requirement already satisfied: typing-extensions>=3.6.6 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorflow) (4.12.2)

Requirement already satisfied: wrapt>=1.11.0 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorflow) (1.16.0)

Requirement already satisfied: grpcio<2.0,>=1.24.3 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorflow) (1.68.0)

Requirement already satisfied: tensorboard<2.19,>=2.18 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorflow) (2.18.0)

Requirement already satisfied: keras>=3.5.0 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorflow) (3.6.0)

Requirement already satisfied: numpy<2.1.0,>=1.26.0 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorflow) (2.0.2)

Requirement already satisfied: h5py>=3.11.0 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorflow) (3.12.1)

Requirement already satisfied: ml-dtypes<0.5.0,>=0.4.0 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorflow) (0.4.1)

Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorflow) (0.37.1)

Requirement already satisfied: wheel<1.0,>=0.23.0 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from astunparse>=1.6.0->tensorflow) (0.45.0)

Requirement already satisfied: rich in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from keras>=3.5.0->tensorflow) (13.9.4)

Requirement already satisfied: namex in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from keras>=3.5.0->tensorflow) (0.0.8)

Requirement already satisfied: optree in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from keras>=3.5.0->tensorflow) (0.13.1)

Requirement already satisfied: charset-normalizer<4,>=2 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from requests<3,>=2.21.0->tensorflow) (3.4.0)

Requirement already satisfied: idna<4,>=2.5 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from requests<3,>=2.21.0->tensorflow) (3.10)

Requirement already satisfied: urllib3<3,>=1.21.1 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from requests<3,>=2.21.0->tensorflow) (2.2.3)

Requirement already satisfied: certifi>=2017.4.17 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from requests<3,>=2.21.0->tensorflow) (2024.8.30)

Requirement already satisfied: markdown>=2.6.8 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorboard<2.19,>=2.18->tensorflow) (3.7)

Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorboard<2.19,>=2.18->tensorflow) (0.7.2)

Requirement already satisfied: werkzeug>=1.0.1 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from tensorboard<2.19,>=2.18->tensorflow) (3.0.1)

Requirement already satisfied: MarkupSafe>=2.1.1 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from werkzeug>=1.0.1->tensorboard<2.19,>=2.18->tensorflow) (2.1.4)

Requirement already satisfied: markdown-it-py>=2.2.0 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from rich->keras>=3.5.0->tensorflow) (3.0.0)

Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from rich->keras>=3.5.0->tensorflow) (2.18.0)

Requirement already satisfied: mdurl~=0.1 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from markdown-it-py>=2.2.0->rich->keras>=3.5.0->tensorflow) (0.1.2)

Requirement already satisfied: wordcloud in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (1.9.4)

Requirement already satisfied: numpy>=1.6.1 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from wordcloud) (2.0.2)

Requirement already satisfied: pillow in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from wordcloud) (11.0.0)

Requirement already satisfied: matplotlib in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from wordcloud) (3.9.2)

Requirement already satisfied: contourpy>=1.0.1 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from matplotlib->wordcloud) (1.3.1)

Requirement already satisfied: cycler>=0.10 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from matplotlib->wordcloud) (0.12.1)

Requirement already satisfied: fonttools>=4.22.0 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from matplotlib->wordcloud) (4.55.0)

Requirement already satisfied: kiwisolver>=1.3.1 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from matplotlib->wordcloud) (1.4.7)

Requirement already satisfied: packaging>=20.0 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from matplotlib->wordcloud) (24.2)

Requirement already satisfied: pyparsing>=2.3.1 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from matplotlib->wordcloud) (3.2.0)

Requirement already satisfied: python-dateutil>=2.7 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from matplotlib->wordcloud) (2.9.0.post0)

Requirement already satisfied: six>=1.5 in /Users/evelynhaskins/.pyenv/versions/3.10.12/lib/python3.10/site-packages (from python-dateutil>=2.7->matplotlib->wordcloud) (1.16.0)

```
In [70]: import os
os.environ['KERAS_BACKEND'] = 'tensorflow'
```

```
In [71]: import os
import numpy as np
import pandas as pd
import tensorflow as tf
import keras_core as keras
import keras_nlp
from sklearn.model_selection import train_test_split
from sklearn.metrics import ConfusionMatrixDisplay, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
import matplotlib.pyplot as plt
import seaborn as sns
from wordcloud import WordCloud
```

```
In [72]: os.environ['KERAS_BACKEND'] = 'tensorflow'

print("TensorFlow version:", tf.__version__)
print("KerasNLP version:", keras_nlp.__version__)
```

TensorFlow version: 2.18.0

KerasNLP version: 0.17.0

Loading the dataset

```
In [73]: base_path = "/Users/evelynhaskins/Downloads/nlp-getting-started"
test_data = os.path.join(base_path, "test.csv")
test = pd.read_csv(test_data)
```

```
train_data = os.path.join(base_path, "train.csv")
train = pd.read_csv(train_data)
```

Evaluating the dataset properties

```
In [74]: print('Training Set Shape = {}'.format(train.shape))
print('Training Set Memory Usage = {:.2f} MB'.format(train.memory_usage().sum() / 1024))
print('Test Set Shape = {}'.format(test.shape))
print('Test Set Memory Usage = {:.2f} MB'.format(test.memory_usage().sum() / 1024))

print(train.head())
print(test.head())

train["length"] = train["text"].apply(lambda x: len(x))
test["length"] = test["text"].apply(lambda x: len(x))

print("Train Length Stats:")
print(train["length"].describe())
print()

print("Test Length Stats:")
print(test["length"].describe())
```

Training Set Shape = (7613, 5)

Training Set Memory Usage = 0.29 MB

Test Set Shape = (3263, 4)

Test Set Memory Usage = 0.10 MB

	id	keyword	location	text	\
0	1	NaN	NaN	Our Deeds are the Reason of this #earthquake M...	
1	4	NaN	NaN	Forest fire near La Ronge Sask. Canada	
2	5	NaN	NaN	All residents asked to 'shelter in place' are ...	
3	6	NaN	NaN	13,000 people receive #wildfires evacuation or...	
4	7	NaN	NaN	Just got sent this photo from Ruby #Alaska as ...	

target

0	1
1	1
2	1
3	1
4	1

	id	keyword	location	text
0	0	NaN	NaN	Just happened a terrible car crash
1	2	NaN	NaN	Heard about #earthquake is different cities, s...
2	3	NaN	NaN	there is a forest fire at spot pond, geese are...
3	9	NaN	NaN	Apocalypse lighting. #Spokane #wildfires
4	11	NaN	NaN	Typhoon Soudelor kills 28 in China and Taiwan

Train Length Stats:

count	7613.000000
mean	101.037436
std	33.781325
min	7.000000
25%	78.000000
50%	107.000000
75%	133.000000
max	157.000000

Name: length, dtype: float64

Test Length Stats:

count	3263.000000
mean	102.108183
std	33.972158
min	5.000000
25%	78.000000
50%	109.000000
75%	134.000000
max	151.000000

Name: length, dtype: float64

Splitting dataset

```
In [75]: BATCH_SIZE = 32
NUM_TRAINING_EXAMPLES = train.shape[0]
TRAIN_SPLIT = 0.8
VAL_SPLIT = 0.2
STEPS_PER_EPOCH = int(NUM_TRAINING_EXAMPLES) * TRAIN_SPLIT // BATCH_SIZE
EPOCHS = 2
AUTO = tf.data.experimental.AUTOTUNE

X = train["text"]
```

```

y = train["target"]
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=VAL_SPLIT,

X_test = test["text"]

print("Training Data Shape:", X_train.shape)
print("Validation Data Shape:", X_val.shape)

```


Training Data Shape: (6090,)
 Validation Data Shape: (1523,)

Analysis of Dataset

```

In [76]: plt.figure(figsize=(10, 6))
sns.histplot(train["length"], bins=30, kde=True, color="blue", label="Train")
sns.histplot(test["length"], bins=30, kde=True, color="orange", label="Test")
plt.title("Distribution of Text Lengths")
plt.xlabel("Text Length")
plt.ylabel("Frequency")
plt.legend()
plt.show()


```



```

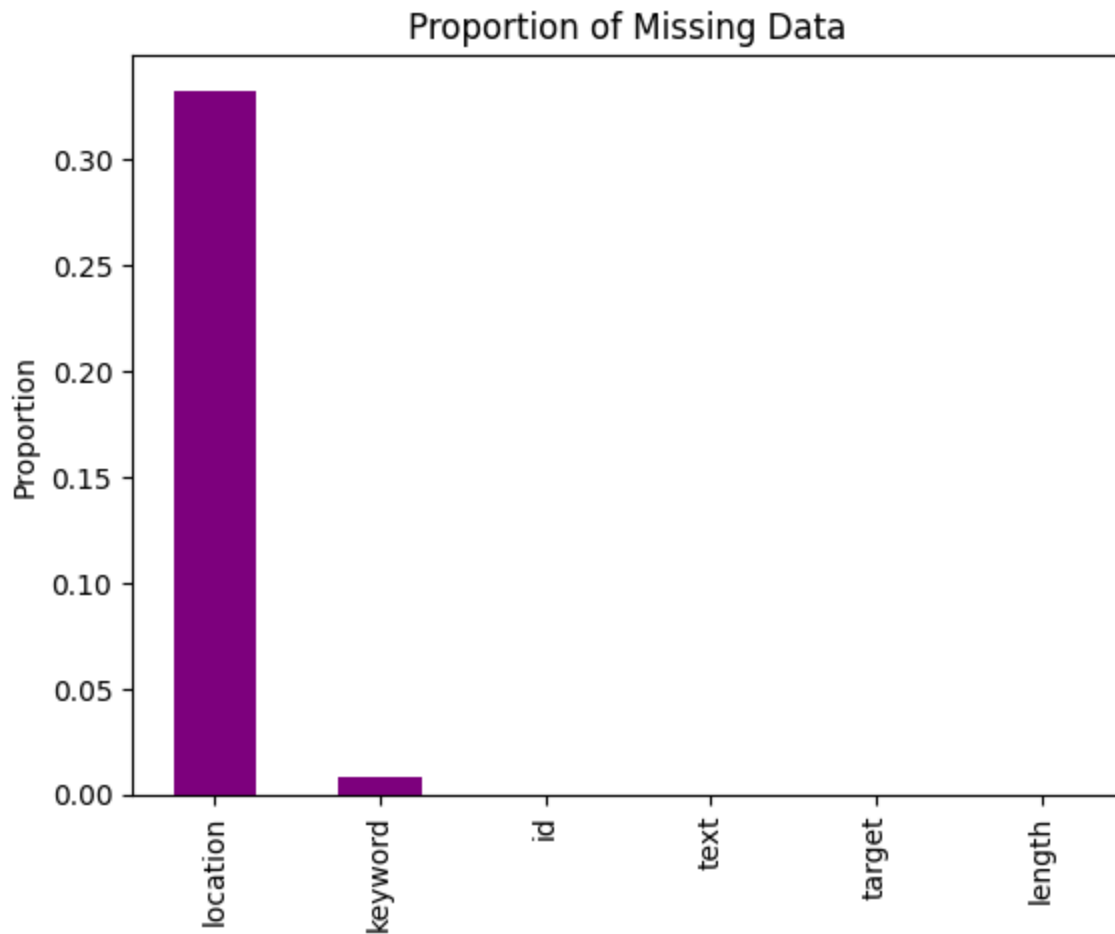
In [77]: train["target"].value_counts().plot(kind="bar", color=["blue", "orange"])
plt.title("Target Class Distribution")
plt.xlabel("Class")
plt.ylabel("Count")
plt.xticks(ticks=[0, 1], labels=["Class 0", "Class 1"], rotation=0)
plt.show()

```



Checking missing values

```
In [78]: missing_data = train.isnull().mean().sort_values(ascending=False)
missing_data.plot(kind="bar", color="purple")
plt.title("Proportion of Missing Data")
plt.ylabel("Proportion")
plt.show()
```

No missing values in areas we are evaluating such as text and target

Checking frequently used words

```
In [79]: text = " ".join(train["text"].fillna("").tolist())
wordcloud = WordCloud(width=800, height=400, background_color="white").generate(text)
plt.figure(figsize=(10, 6))
plt.imshow(wordcloud, interpolation="bilinear")
plt.axis("off")
plt.title("Most Common Words in Text")
plt.show()
```



Interesting way to see common words trending in the tweets, not overly relevant in this evaluation but fun to look at

```
In [80]: preset = "distil_bert_base_en_uncased"
preprocessor = keras_nlp.models.DistilBertPreprocessor.from_preset(
    preset, sequence_length=160, name="preprocessor_4_tweets"
)
classifier = keras_nlp.models.DistilBertClassifier.from_preset(
    preset, preprocessor=preprocessor, num_classes=2
)
classifier.summary()
```

Preprocessor: "preprocessor_4_tweets"

Layer (type)	
distil_bert_tokenizer (DistilBertTokenizer)	

Model: "distil_bert_text_classifier_4"

Layer (type)	Output Shape	
padding_mask (InputLayer)	(None, None)	
token_ids (InputLayer)	(None, None)	
distil_bert_backbone (DistilBertBackbone)	(None, None, 768)	66,
get_item_4 (GetItem)	(None, 768)	
pooled_dense (Dense)	(None, 768)	
output_dropout (Dropout)	(None, 768)	
logits (Dense)	(None, 2)	

Total params: 66,955,010 (255.41 MB)

Trainable params: 66,955,010 (255.41 MB)

Non-trainable params: 0 (0.00 B)

Creating a Neural Network to predict target for tweets to determine whether a person's words are actually announcing a disaster.

Convert the data into TensorFlow Datasets for efficient batching

```
In [81]: def prepare_dataset(X, y, batch_size, is_training=True):
dataset = tf.data.Dataset.from_tensor_slices((X, y))
if is_training:
dataset = dataset.shuffle(len(X))
dataset = dataset.batch(batch_size).prefetch(AUTO)
return dataset
```

These lines of code create TensorFlow datasets for efficient training and validation by converting the raw data into a format that is optimized for use with TensorFlow models.

```
In [82]: train_dataset = prepare_dataset(X_train, y_train, BATCH_SIZE, is_training=True)
val_dataset = prepare_dataset(X_val, y_val, BATCH_SIZE, is_training=False)
```

```
In [83]: # Ensure the datasets repeat for continuous training across epochs
train_dataset = train_dataset.repeat()
val_dataset = val_dataset.repeat()
```

Compiling the data: Use the Adam optimizer to adjust the weights during backpropagation. Measure the model's performance using Sparse Categorical Crossentropy as the loss function. Track and display the model's accuracy during training and validation.

```
In [92]: classifier.compile(
optimizer=tf.keras.optimizers.Adam(learning_rate=5e-5),
loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
```

```
metrics=['accuracy']
)
```

Reduced size just for this project so I didn't have to wait for it to run all day, wouldn't recommend this

```
In [97]: train_dataset = train_dataset.prefetch(buffer_size=tf.data.experimental.AUTOTUNE)
val_dataset = val_dataset.prefetch(buffer_size=tf.data.experimental.AUTOTUNE)
```

Training the neural network

```
In [98]: # Train the model
history = classifier.fit(
    train_dataset.take(100), # Use a small subset
    validation_data=val_dataset.take(10),
    epochs=10,
    steps_per_epoch=10,
    validation_steps=2
)
```

```
Epoch 1/10
10/10 _____ 36s 4s/step - accuracy: 0.8379 - loss: 0.4081 - val_accuracy: 0.7500 - val_loss: 0.4842
Epoch 2/10
10/10 _____ 36s 4s/step - accuracy: 0.8219 - loss: 0.4003 - val_accuracy: 0.8125 - val_loss: 0.4688
Epoch 3/10
10/10 _____ 36s 4s/step - accuracy: 0.8684 - loss: 0.3452 - val_accuracy: 0.7812 - val_loss: 0.4704
Epoch 4/10
10/10 _____ 36s 4s/step - accuracy: 0.7973 - loss: 0.4822 - val_accuracy: 0.8125 - val_loss: 0.4202
Epoch 5/10
10/10 _____ 35s 4s/step - accuracy: 0.8269 - loss: 0.4325 - val_accuracy: 0.8281 - val_loss: 0.4134
Epoch 6/10
10/10 _____ 35s 4s/step - accuracy: 0.8271 - loss: 0.4171 - val_accuracy: 0.8438 - val_loss: 0.3692
Epoch 7/10
10/10 _____ 35s 4s/step - accuracy: 0.8904 - loss: 0.3044 - val_accuracy: 0.8125 - val_loss: 0.3836
Epoch 8/10
10/10 _____ 36s 4s/step - accuracy: 0.8683 - loss: 0.3672 - val_accuracy: 0.8438 - val_loss: 0.3901
Epoch 9/10
10/10 _____ 36s 4s/step - accuracy: 0.8425 - loss: 0.3701 - val_accuracy: 0.8125 - val_loss: 0.4279
Epoch 10/10
10/10 _____ 36s 4s/step - accuracy: 0.8235 - loss: 0.4184 - val_accuracy: 0.8281 - val_loss: 0.4517
```

Why I Chose DistilBERT for This Architect

For this text classification problem, I chose **DistilBERT**, a transformer-based architecture, because it provides an efficient and effective solution for processing natural language. DistilBERT is a smaller, faster, and lighter version of the well-known **BERT (Bidirectional Encoder Representations from Transformers)**, which is built on the **transformer architecture** introduced in the seminal paper "*Attention is All You Need*".

Key reasons for choosing DistilBERT:

1. Transformer-Based Architecture:

DistilBERT retains the core components of the transformer architecture, including the **self-attention mechanism**, which allows it to capture contextual relationships between words in a sequence effectively. This is critical for tasks where understanding the meaning of text depends on context.

2. Efficiency Without Sacrificing Performance:

DistilBERT is designed to be faster and more resource-efficient than the original BERT. It achieves this by:

- Reducing the number of layers (6 in DistilBERT vs. 12 in BERT base).
- Using knowledge distillation to transfer knowledge from a larger BERT model while maintaining comparable accuracy.

This makes it well-suited for my dataset, as it balances computational cost and performance, enabling faster training and inference.

3. Pre-Trained Model:

By using a pre-trained version of DistilBERT, I leverage the knowledge it has already learned from large corpora, such as Wikipedia and BookCorpus. This helps achieve better results with limited training data compared to training a model from scratch.

4. Flexibility for Text Classification:

DistilBERT's pre-trained model can easily be extended with a classification head, which is specifically designed for binary classification tasks like mine. Using the `keras_nlp.models.DistilBertClassifier`, I benefit from an end-to-end pipeline tailored for this problem.

Evaluating Accuracy

```
In [99]: accuracy = history.history['accuracy']
val_accuracy = history.history['val_accuracy']

epochs = range(1, len(accuracy) + 1)


# Plot the accuracy
plt.figure(figsize=(8, 6))
plt.plot(epochs, accuracy, 'bo-', label='Training Accuracy')
plt.plot(epochs, val_accuracy, 'ro-', label='Validation Accuracy')
```

```
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)
plt.show()
```



```
In [100... loss = history.history['loss']
val_loss = history.history['val_loss']

# Plot the loss
plt.figure(figsize=(8, 6))
plt.plot(epochs, loss, 'bo-', label='Training Loss')
plt.plot(epochs, val_loss, 'ro-', label='Validation Loss')
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.show()
```



```
In [101]... def displayConfusionMatrix(y_true, y_pred, dataset):
    disp = ConfusionMatrixDisplay.from_predictions(
        y_true,
        np.argmax(y_pred, axis=1),
        display_labels=["Not Disaster", "Disaster"],
        cmap=plt.cm.Blues
    )


    tn, fp, fn, tp = confusion_matrix(y_true, np.argmax(y_pred, axis=1)).ravel()
    f1_score = tp / (tp + ((fn + fp) / 2))

    disp.ax_.set_title("Confusion Matrix on " + dataset + " Dataset -- F1 Score")
```

```
In [102]... y_pred_train = classifier.predict(X_train)

    displayConfusionMatrix(y_train, y_pred_train, "Training")
```

191/191 ————— 220s 1s/step



The **confusion matrix** above is used to evaluate the performance of the classification model, and it helps you understand how well the model is distinguishing between different Disaster and Not Disaster.

Structure of the Confusion Matrix

For a binary classification problem, a confusion matrix looks like this:

	Predicted Positive (1)	Predicted Negative (0)
Actual Positive (1)	True Positive (TP)	False Negative (FN)
Actual Negative (0)	False Positive (FP)	True Negative (TN)

Where:

- **True Positive (TP):** Correctly predicted Disaster Tweets.
- **False Positive (FP):** Incorrectly predicted Disaster Tweets (weren't actually disaster tweets).
- **False Negative (FN):** Incorrectly predicted as Not Disaster (were Disaster Tweets).
- **True Negative (TN):** Correctly predicted Not Disaster Tweets.

In []: